

## RCOM Project 2

Gonçalo Miranda (up202108773)  
Sophie Large (up202303141)

### Table of Contents

#### RCOM Project 2

##### Table of Contents

##### Introduction

##### Part 1 - Development of the Download Application (app.c and app.h)

##### Architecture of the Application

##### Part 2 - Experiences

##### Experience 1 - Configure an IP Network

##### Step By Step

##### Questions

##### Experience 2 - Implement two bridges in a switch

##### Step By Step

##### Questions

##### Experience 3 - Configure a router in Linux

##### Step By Step

##### Questions

##### Experience 4 - Configure a Commercial Router and Implement NAT

##### Step By Step

##### Questions

##### Experience 5 - DNS

##### Step by Step

##### Questions

##### Experience 6 - TCP Connections

##### Step By Step

##### Questions

##### Annexes

##### 1. Code

##### 2. Wireshark Captures

# Introduction

The goal of this project was to set up a small network and prove that it is correctly setup using a download application, which was also developed for this project.

## Part 1 - Development of the Download Application (app.c and app.h)

This part of the project consists in the development of a small app that downloads files using the FTP protocol. The app works when provided with an URL such as  
`ftp://[<username>[:<password>]@]<host>[:<port>]/<path>.`

Example: `ftp://ftp.up.pt/pub/hello.txt`

During the development of the app part of the project, we had the chance to understand Unix utilities to handle URLs and to communicate with sockets, the FTP protocol and the concept of DNS.

## Architecture of the Application

1. Parsing the URL string given as the argument to the program to extract its fragments, such as the username, password, host, port, and path.
2. Connecting to the host on the specified port (being 21 the default port, the standard FTP control port).
3. Entering passive mode and receiving the host and port of the data connection.
4. Connecting to the data connection host and port.
5. Notifying the server that it should start transferring the desired file.
6. Reading the incoming data from the server and writing it locally to the download. file, being the extension of the last part of the path (last string after the last "/").
7. Closing the connections and terminating the app.

TCP sockets mediate the network communication. By making use of the URL parser we designed, via a state machine and regular expressions, and `getaddrinfo()`, the FTP server's IP and control port are obtained (in case there's no port specified, 21 is the value assumed). A TCP connection is established and the program starts sending commands through the control connection, which will be responded to by the server. Those responses carry response codes that are properly extracted and checked to understand if the program is working as intended. If any of the response codes sent by the server doesn't match the expected code, the program is terminated with exit code -1. In case there's no username specified by the user, "anonymous" is assumed as the username, since it's the default value we decided to provide. After that, passive mode is entered and, if its successfully entered, the server responds with code 227 and the host's IP and port. After establishing a data connection, the `retr` command is sent and, after that, the buffers of data from the data connection begin getting read into a file named "download.<extension>", where the extension is the last string after the last "/" in the file path provided by the user. Finally, the connection is closed and the program terminates.

## Part 2 - Experiences

### Experience 1 - Configure an IP Network

#### Step By Step

Note that, in this experience, the workbench used was workbench 4.

1. Connect E0 on TuxY3 and TuxY4 on any of the switch ports.
2. Configure TuxY3 and TuxY4, by performing the following commands:
  - TuxY3

```
$ ifconfig eth0 up
$ ifconfig eth0 172.16.Y0.1/24
```
  - TuxY4

```
$ ifconfig eth0 up
$ ifconfig eth0 172.16.Y0.254/24
```
3. Now we can check both the IP and MAC addresses of both computers:
  - TuxY3:
    - IP Address: 172.16.Y0.1/24
    - MAC Address: 00:21:5a:61:2f:d4
  - TuxY4:
    - IP Address: 172.16.Y0.254/24
    - MAC Address: 00:21:5a:5a:7b:ea
4. We now can test the connectivity between the computers by pinging each other (TuxY3 pings TuxY4 and vice-versa).
  - TuxY3

```
$ ping 172.16.Y0.254 -c 10
```

```
root@tux43:~# ping 172.16.40.254 -c 10
PING 172.16.40.254 (172.16.40.254) 56(84) bytes of data:
64 bytes from 172.16.40.254: icmp_seq=1 ttl=64 time=0.150 ms
64 bytes from 172.16.40.254: icmp_seq=2 ttl=64 time=0.139 ms
64 bytes from 172.16.40.254: icmp_seq=3 ttl=64 time=0.139 ms
64 bytes from 172.16.40.254: icmp_seq=4 ttl=64 time=0.183 ms
64 bytes from 172.16.40.254: icmp_seq=5 ttl=64 time=0.134 ms
64 bytes from 172.16.40.254: icmp_seq=6 ttl=64 time=0.139 ms
64 bytes from 172.16.40.254: icmp_seq=7 ttl=64 time=0.151 ms
64 bytes from 172.16.40.254: icmp_seq=8 ttl=64 time=0.146 ms
64 bytes from 172.16.40.254: icmp_seq=9 ttl=64 time=0.187 ms
64 bytes from 172.16.40.254: icmp_seq=10 ttl=64 time=0.136 ms

--- 172.16.40.254 ping statistics ---
10 packets transmitted, 10 received, 0% packet loss, time 121ms
rtt min/avg/max/mdev = 0.134/0.150/0.187/0.021 ms
```

- TuxY4:

```
$ ping 172.16.Y0.1 -c 10
```

```
root@tux44:~# ping 172.16.40.1 -c 10
PING 172.16.40.1 (172.16.40.1) 56(84) bytes of data.
64 bytes from 172.16.40.1: icmp_seq=1 ttl=64 time=0.150 ms
64 bytes from 172.16.40.1: icmp_seq=2 ttl=64 time=0.145 ms
64 bytes from 172.16.40.1: icmp_seq=3 ttl=64 time=0.132 ms
64 bytes from 172.16.40.1: icmp_seq=4 ttl=64 time=0.138 ms
64 bytes from 172.16.40.1: icmp_seq=5 ttl=64 time=0.182 ms
64 bytes from 172.16.40.1: icmp_seq=6 ttl=64 time=0.155 ms
64 bytes from 172.16.40.1: icmp_seq=7 ttl=64 time=0.150 ms
64 bytes from 172.16.40.1: icmp_seq=8 ttl=64 time=0.140 ms
64 bytes from 172.16.40.1: icmp_seq=9 ttl=64 time=0.134 ms
64 bytes from 172.16.40.1: icmp_seq=10 ttl=64 time=0.180 ms

--- 172.16.40.1 ping statistics ---
10 packets transmitted, 10 received, 0% packet loss, time 224ms
rtt min/avg/max/mdev = 0.132/0.150/0.182/0.021 ms
```

The connection is working in case all the packets are correctly received.

5. Now, we can check the forwarding and the ARP tables. We're doing it on TuxY3.

In:

```
$ route -n
```

```
out:


| Destination | Gateway       | Genmask       | Flags | Metric | Ref | Use | Iface |
|-------------|---------------|---------------|-------|--------|-----|-----|-------|
| 0.0.0.0     | 172.16.40.254 | 0.0.0.0       | UG    | 0      | 0   | 0   | eth0  |
| 172.16.40.0 | 172.16.40.254 | 255.255.255.0 | UG    | 0      | 0   | 0   | eth0  |
| 172.16.40.0 | 0.0.0.0       | 255.255.255.0 | U     | 0      | 0   | 0   | eth0  |
| 172.16.41.0 | 172.16.40.254 | 255.255.255.0 | UG    | 0      | 0   | 0   | eth0  |


```

In:

```
$ arp -a
```

out:

```
? (172.16.Y0.254) at 00:21:5a:5a:7b:ea [ether] on eth0
```

6. We are now deleting the arp table entries in TuxY3. In:

```
$ arp -d 172.16.Y0.254
```

To test if the deletion occurred correctly, we can perform

In:

```
$ arp -a
```

out:

No entries on the table

7. Now, we can start Wireshark in TuxY3.eth0 and start capturing packets.

8. After that, we can ping TuxY4 from TuxY3 10 times.

```
$ ping 172.16.40.254 -c 10
```

9. Finally, we can stop capturing the packets.

10. When saving the logs and analyzing them, we've got the following:

32	172.16.40.254	48.154358026	172.16.40.1	ICMP	98	Echo (ping) request	id=0x1f3d, seq=1/256, ttl=64	(reply in 33)
33	172.16.40.1	48.154480248	172.16.40.254	ICMP	98	Echo (ping) reply	id=0x1f3d, seq=1/256, ttl=64	(request in 32)
34	172.16.40.254	49.157306013	172.16.40.1	ICMP	98	Echo (ping) request	id=0x1f3d, seq=2/512, ttl=64	(reply in 35)
35	172.16.40.1	49.157464832	172.16.40.254	ICMP	98	Echo (ping) reply	id=0x1f3d, seq=2/512, ttl=64	(request in 34)
37	172.16.40.254	50.181305727	172.16.40.1	ICMP	98	Echo (ping) request	id=0x1f3d, seq=3/768, ttl=64	(reply in 38)
38	172.16.40.1	50.181434375	172.16.40.254	ICMP	98	Echo (ping) reply	id=0x1f3d, seq=3/768, ttl=64	(request in 37)
39	172.16.40.254	51.205307955	172.16.40.1	ICMP	98	Echo (ping) request	id=0x1f3d, seq=4/1024, ttl=64	(reply in 40)
40	172.16.40.1	51.205431225	172.16.40.254	ICMP	98	Echo (ping) reply	id=0x1f3d, seq=4/1024, ttl=64	(request in 39)
42	172.16.40.254	52.229303827	172.16.40.1	ICMP	98	Echo (ping) request	id=0x1f3d, seq=5/1280, ttl=64	(reply in 43)
43	172.16.40.1	52.229426818	172.16.40.254	ICMP	98	Echo (ping) reply	id=0x1f3d, seq=5/1280, ttl=64	(request in 42)
44	172.16.40.254	53.253302005	172.16.40.1	ICMP	98	Echo (ping) request	id=0x1f3d, seq=6/1536, ttl=64	(reply in 45)
45	172.16.40.1	53.253429186	172.16.40.254	ICMP	98	Echo (ping) reply	id=0x1f3d, seq=6/1536, ttl=64	(request in 44)
49	172.16.40.254	54.277302347	172.16.40.1	ICMP	98	Echo (ping) request	id=0x1f3d, seq=7/1792, ttl=64	(reply in 50)
50	172.16.40.1	54.277445452	172.16.40.254	ICMP	98	Echo (ping) reply	id=0x1f3d, seq=7/1792, ttl=64	(request in 49)
51	172.16.40.254	55.301306182	172.16.40.1	ICMP	98	Echo (ping) request	id=0x1f3d, seq=8/2048, ttl=64	(reply in 52)
52	172.16.40.1	55.301435388	172.16.40.254	ICMP	98	Echo (ping) reply	id=0x1f3d, seq=8/2048, ttl=64	(request in 51)
54	172.16.40.254	56.325301496	172.16.40.1	ICMP	98	Echo (ping) request	id=0x1f3d, seq=9/2304, ttl=64	(reply in 55)
55	172.16.40.1	56.325424905	172.16.40.254	ICMP	98	Echo (ping) reply	id=0x1f3d, seq=9/2304, ttl=64	(request in 54)
56	172.16.40.254	57.349303165	172.16.40.1	ICMP	98	Echo (ping) request	id=0x1f3d, seq=10/2560, ttl=64	(reply in 57)
57	172.16.40.1	57.349432441	172.16.40.254	ICMP	98	Echo (ping) reply	id=0x1f3d, seq=10/2560, ttl=64	(request in 56)
75	ff02::2	90.245279725	fe80::221:5a...	ICMPv6	70	Router Solicitation from 00:21:5a:61:2f:d4		

## Questions

- What are the ARP packets and what are they used for?  
The ARP (*Address Resolution Protocol*) packets are used to establish a connection between an IP and a MAC address.
- What are the MAC and IP addresses of ARP packets and why?  
The ARP packets carry 2 IP addresses: the destiny machine's IP address (172.16.Y0.254, TuxY4) the the source machine's IP address (172.16.Y0.01, TuxY3). TuxY4 then proceeds to send another ARP packet to TuxY3 that contains its MAC address (00:21:5a:5a:7b:ea).
- What packets does the ping command generate?  
While the destination's MAC address isn't received, *ping* generates ARP packets. After connecting the IP address to the corresponding MAC address, it generates ICMP (*Internet Control Message Protocol*) packets.
- What are the MAC and IP addresses of the ping packets?  
The IP and MAC addresses used in ICMP are the ones from the machines in use - in this case, TuxY3 and TuxY4.
- How to determine if a receiving Ethernet frame is ARP, IP, ICMP?  
Each frame's protocol designation can be checked in the "Protocol" column of the Wireshark capture. The protocol is usually indicated in the first 2 bytes of the frame, in its header.
- How to determine the length of a receiving frame?  
Wireshark has a column that specifies the frame size, in bytes. Usually that is also present in the frame's header.
- What is the loopback interface and why is it important?  
It's a virtual interface that's always reachable as long as one of the switch's interfaces are on. Because of this, it's possible to periodically check if the network's connections are properly configured.

## Experience 2 - Implement two bridges in a switch

### Step By Step

Note that, in this experience, the workbench used was workbench 4.

1. Connect E0 on TuxY2 to any of the switch ports. After that, configure TuxY2, by performing the following commands:
 

```
$ ifconfig eth0 up
$ ifconfig eth0 172.16.Y1.1/24
```
2. Now, we can create two bridges in the switch, using the following commands on GTKTerm.
  - BridgeY0:
 

```
/interface bridge add name=bridgeY0
```
  - BridgeY1:
 

```
/interface bridge add name=bridgeY1
```
3. Afterwards, we can remove the ports where TuxY2, TuxY3, and TuxY4 are connected from the default bridge and add them to the corresponding ports to bridgeY0 and bridgeY1.
  - Removing the original ports:
 

```
/interface bridge port remove [find interface=ether1]
/interface bridge port remove [find interface=ether2]
/interface bridge port remove [find interface=ether3]
```
  - Adding the current ports:
 

```
/interface bridge port add bridge=bridgeY0 interface=ether1
/interface bridge port add bridge=bridgeY0 interface=ether2
/interface bridge port add bridge=bridgeY1 interface=ether3
```
4. Now, we should switch to TuxY3 and start monitoring the packets.
5. Ping TuxY4 and TuxY2 from TuxY3.
 

```
In:
$ ping 172.16.Y0.254 -c 10
Out:
In TuxY4, all 10 packets have been received successfully!

In:
$ ping 172.16.Y1.1 -c 10
Out:
In TuxY2 -> connect: Network is unreachable
```
6. Afterwards, stop the capture.
7. Proceed by starting the capture of eth0 in TuxY2, TuxY3 and TuxY4.
8. From TuxY3, ping 172.16.Y0.255.
 

```
$ ping 172.16.Y0.255 -c 10
```
9. Stop the captures.
10. Restart the previous captures.
11. From TuxY2, ping 172.16.Y1.255.
12. Enable ICMP.
13. Enable IP forwarding.

## Questions

1. How to configure bridgeY0?  
In order to connect TuxY3 and TuxY4, bridgeY0 was configured, creating a subnet. To configure the bridge properly, we delete previous configurations and default connections when the machines first connected and configure new ports that connect to each

machine. We can do this by using the commands stated in the previous step-by-step guide.

2. How many broadcast domains are there? How can you conclude it from the logs?  
Since we implemented two bridges (bridgeY0 and bridgeY1), there are two broadcast domains. We can conclude this because, after analyzing the Wireshark logs, TuxY3 got a response from TuxY4 but not from TuxY2, since they're on different bridges and, therefore, different broadcast domains.

## Experience 3 - Configure a router in Linux

### Step By Step

Note that, in this experience, the workbench used was workbench 5.

1. We need to configure TuxY4 as a router, by:
  - Connect TuxY4's eth1 to a switch port.
  - After connecting the cables, config the port itself.  
`$ ifconfig eth1 172.16.Y1.253/24`
  - Connect TuxY2 to the switch console and add it to bridgeY1.  
In GTKTerm:  
`/interface bridge port remove [find interface=etherXX]`  
`/interface bridge port add interface=etherXX bridge=bridgeY1`  
, being XX the switch's port number.
  - Enable the IP forwarding on TuxY4.  
`$ echo 1 > /proc/sys/net/ipv4/ip_forward`
  - Disable ICMP echo-ignore broadcast on TuxY4.  
`$ echo 1 > /proc/sys/net/ipv4/icmp_echo_ignore_broadcasts`
2. Now, check the MAC and IP addresses in TuxY4's eth0 and eth1.
  - in:  
`$ ifconfig`
  - out:
    - i. eth0:  
MAC Address: 00:21:5a:c3:78:70 IP Address: 172.16.Y0.254
    - ii. eth1:  
MAC Address: 00:c0:df:08:d5:b0 IP Address: 172.16.Y1.253
3. Reconfigure TuxY3 and TuxY2 so that they can reach each other.
  - TuxY2:  
`$ route add -net 172.16.50.0/24 gw 172.16.51.253`
  - TuxY3:  
`$ route add -net 172.16.Y1.0/24 gw 172.16.Y0.254`
4. Now, check the routes available in each machine.  
`$ route -n`
5. Start capturing packets with Wireshark.

6. From TuxY3, ping the other network interfaces and check their connectivity.  
in:

```
$ ping 172.16.Y0.254 -c 10
```

```
$ ping 172.16.Y1.253 -c 10
```

```
$ ping 172.16.Y1.1 -c 10
```

out:

All 10 packets for each ping call received correctly in every machine.
7. Stop the log capture and check the logs.
8. Now, start a Wireshark capture on TuxY4's eth0 and eth1.
9. In each machine, clean the ARP tables.
  - o TuxY2:

```
$ arp -d 172.16.Y1.253
```
  - o TuxY3:

```
$ arp -d 172.16.Y0.254
```
  - o TuxY4:

```
$ arp -d 172.16.Y0.1 $ arp -d 172.16.Y1.1
```
10. In TuxY3, ping TuxY2.

```
$ ping 172.16.Y1.1 c 10
```
11. Stop both TuxY4 Wireshark captures.

## Questions

1. What are the commands required to configure this experience?  
The necessary commands to configure the experience are the ones displayed in the step-by-step guide presented before.
2. What routes are there in the tuxes? What are their meaning?  
There is a route on TuxY2 and another in TuxY3. Both those routes have TuxY4 as the gateway, since TuxY4 is the only one common to the created bridges (it's a router).
3. What information does an entry of the forwarding table contain?  
Each entry of the forwarding table has the destiny address and its gateway.
4. What ARP messages, and associated MAC addresses, are observed and why?  
When pinging TuxY2 from TuxY3, the ARP packages carry the IP and MAC addresses of TuxY3 and TuxY4 instead of the TuxY2 as the destiny address. This happens because of the routing we created: here, TuxY2's address is not known by TuxY3. The latter only knows the TuxY4's address (gateway), that will later lead to TuxY2.

40 50.795323415	HewlettP_61:2d:72	HewlettP_c3:78:70	ARP	42 Who has 172.16.50.254? Tell 172.16.50.1
41 50.795457161	HewlettP_c3:78:70	HewlettP_61:2d:72	ARP	60 172.16.50.254 is at 00:21:5a:c3:78:70
42 50.799918204	HewlettP_c3:78:70	HewlettP_61:2d:72	ARP	60 Who has 172.16.50.1? Tell 172.16.50.254
43 50.799926235	HewlettP_61:2d:72	HewlettP_c3:78:70	ARP	42 172.16.50.1 is at 00:21:5a:61:2d:72

5. What ICMP packets are observed and why?  
The ICMP packets carry the source (TuxY3) and destiny's IP addresses (TuxY2). The fact that those two are the IP's carried means the network is correctly configured.

97 172.473995894	HewlettP_61:2d:72	Broadcast	ARP	60 Who has 172.16.50.254? Tell 172.16.50.1
98 172.474016498	HewlettP_c3:78:70	HewlettP_61:2d:72	ARP	42 172.16.50.254 is at 00:21:5a:c3:78:70
99 172.474128872	172.16.50.1	172.16.51.1	ICMP	98 Echo (ping) request id=0x1532, seq=1/256, ttl=64 (reply in 100)
100 172.474386307	172.16.51.1	172.16.50.1	ICMP	98 Echo (ping) reply id=0x1532, seq=1/256, ttl=63 (request in 99)
101 173.500297074	172.16.50.1	172.16.51.1	ICMP	98 Echo (ping) request id=0x1532, seq=2/512, ttl=64 (reply in 102)
102 173.500449049	172.16.51.1	172.16.50.1	ICMP	98 Echo (ping) reply id=0x1532, seq=2/512, ttl=63 (request in 101)

6. What are the IP and MAC addresses associated to ICMP packets and why?  
Each ICMP packet observed after the ping from TuxY3 to TuxY2, carry TuxY3's IP address as the source IP address and the TuxY3's IP address as the destiny address,



however, the ICMP packets carry TuxY4's MAC address, due to the fact that this machine is the one connecting the two bridges.

## Experience 4 - Configure a Commercial Router and Implement NAT

### Step By Step

Note that, in this experience, the workbench used was workbench 5.

1. Connect the router's eth1 port to the port Y1 of the server rack in the lab - NAT is, by default, enabled.
2. Connect the router's eth2 port to the switch.
3. Delete the default ports of the switch's etherXX (XX represents the port number on which the cable is connected) and connect it to bridgeY1.
  - o Using GTKTerm, in the switch's console, run:  

```
/interface bridge port remove [find interface=etherXX]  
/interface bridge port add bridge=bridgeY1 interface=etherXX
```
4. Switch the cable that was connected to the switch to the router.
5. Reset the router's settings and configure its IP address.
  - o Using GTKTerm, switch the baudrate to 115200 and the serial port to /dev/ttyS0. After that, run:  

```
/system reset-configuration
```
  - o Login with username "admin" and no password.
  - o Run:  

```
/ip address add address=172.16.1.Y9/24 interface=ether1 /ip  
address add address=172.16.Y1.254/24 interface=ether2
```
6. Configure all the computer's default routes, as well as the router's.
  - o TuxY2:  

```
$ route add default gw 172.16.Y1.254
```
  - o TuxY3:  

```
$ route add default gw 172.16.Y0.254
```
  - o TuxY4:  

```
$ route add default gw 172.16.Y1.254
```
  - o Using GTKTerm on the router's console, run: 

```
/ip route add  
dst-address=172.16.Y0.0/24 gateway=172.16.Y1.253  
/ip route add dst-address=0.0.0.0/0 gateway=172.16.1.254
```
7. Start a Wireshark capture on TuxY3 to check if it can ping all the network interfaces of TuxY2, TuxY4 and the router.
  - In:  

```
$ ping 172.16.Y0.254 -c 10  
$ ping 172.16.Y1.1 -c 10  
$ ping 172.16.Y1.254 -c 10
```
  - Out:  
All 10 packets for each ping call received correctly in every machine.
8. On TuxY2, disable accept\_redirects.
  - ```
$ sysctl net.ipv4.conf.eth0.accept_redirects=0  
$ sysctl net.ipv4.conf.all.accept_redirects=0
```

9. Remove the route that connects TuxY2 to TuxY4.

```
$ route del -net 172.16.Y0.0 gw 172.16.Y1.253 netmask
255.255.255.0
```

10. On TuxY2, start a Wireshark capture and ping TuxY3. In:

In:

```
$ ping 172.16.Y0.1
```

Out:

All 10 packets of the ping call received correctly (connection was using the proper router instead of TuxY4 as a router).

11. Now, let's check the connection's route. In:

In:

```
$ traceroute -n 172.16.Y0.1
```

Out:

```
traceroute to 172.16.Y0.1 (172.16.Y0.1), 30 hops max, 60 byte packets
1 172.16.Y1.254 (172.16.Y1.254) 0.200 ms 0.204 ms 0.224 ms
2 172.16.Y1.253 (172.16.Y1.253) 0.354 ms 0.345 ms 0.344 ms
3 tuxY1 (172.16.Y0.1) 0.596 ms 0.587 ms 0.575 ms
```

12. Recreate the route that connects TuxY2 to TuxY4.

```
$ route add -net 172.16.Y0.0/24 gw 172.16.Y1.253
```

13. Trace the route from TuxY2 to TuxY3.

In:

```
$ traceroute -n 172.16.Y0.1
```

Out:

```
traceroute to 172.16.Y0.1 (172.16.Y0.1), 30 hops max, 60 byte packets
1 172.16.Y1.253 (172.16.Y1.253) 0.196 ms 0.180 ms 0.164 ms
2 tuxY1 (172.16.Y0.1) 0.414 ms 0.401 ms 0.375 ms
```

14. Reactivate accept\_redirects.

```
$ sysctl net.ipv4.conf.eth0.accept_redirects=1
$ sysctl net.ipv4.conf.all.accept_redirects=1
```

15. From TuxY3, ping the Lab's router.

```
$ ping 172.1.1.254
```

16. Deactivate the router's NAT. On GTKTerm, on the router's console:

```
/ip firewall nat disable 0
```

On TuxY3, ping the Lab's router again.

In:

```
$ ping 172.16.1.254
```

Out:

No connection.

17. Reactivate the router's NAT.

```
/ip firewall nat enable 0
```

## Questions

1. How to configure a static route in a commercial router?  
We start by resetting its settings, add it to the corresponding bridge and give it its own external and internal IP address.
2. What are the paths followed by the packets in the experiments carried out and why?  
In the first part of the experience, where TuxY2 wasn't connected to TuxY4, the data packets were redirected (ICMP redirect) through the implemented router to the destination IP address.  
In the second part of the experience, there were no redirects, since the shortest connection in the network was available.
3. How to configure NAT in a commercial router? Using the GTK Terminal command `/ip firewall nat enable 0` in the router's terminal.
4. What does NAT do? NAT (Network Address Translation) translates addresses from the local network to a single public address and vice-versa. Because of this, when a packet is sent to an external network, the package is sent with the public address as the source. When the destiny machine sends back a response, it sends that response to the public address, which is then re-translated again to the local address that sent the packet in the first place. This is useful, since we can reduce the number of public addresses used on a network.

## Experience 5 - DNS

### Step by Step

Note that, in this experience, the workbench used was workbench 4 and in room I320.

1. Configure DNS in all machines.
  - o TuxY2:  
`$ nano /etc/resolv.conf` and add the line `nameserver 172.16.2.1`
  - o TuxY3:  
`$ nano /etc/resolv.conf` and add the line `nameserver 172.16.2.1`
  - o TuxY4:  
`$ nano /etc/resolv.conf` and add the line `nameserver 172.16.2.1`
2. In all 3 machines, ping website names to check if they can be used as host.  
In:  
`$ ping google.com -c 10`  
Out:  
All 10 packets correctly transmitted and received from every machine.

## Questions

1. How to configure the DNS service at a host?  
DNS is configured by adding the line "nameserver " to the end of the file `/etc/resolv.conf` of all connected machines.
2. What packets are exchanged by DNS and what information is transported? The first few packets transported are DNS packets, so that the router can identify and translate the destination IP address.

# Experience 6 - TCP Connections

## Step By Step

Note that, in this experience, the workbench used was workbench 4.

1. Compile the download application in TuxY3.
2. In TuxY3, start a Wireshark capture and execute the application.
3. Check if the file has correctly arrived and stop the Wireshark capture.
4. Check the Wireshark log and analyze the data packets exchanged.
5. Start the Wireshark capture again and restart the download in TuxY3. After it started, switch to TuxY2 and start a download there too.
6. Stop the capture and analyze the logs.

## Questions

1. How many TCP connections are opened by your FTP application?  
The application has 2 TCP connections, one to send commands and another to receive the file.
2. In what connection is transported the FTP control information?  
The control information is transported through the connection that has been setup to send commands to the server and fetch its responses.
3. What are the phases of a TCP connection?  
A TCP connection is divided in 3 main phases:
  - Connection Establishment (Three-Way Handshake), where the client sends a TCP segment to the server with the **SYN** flag. That is followed by a response from the server, which is a TCP segment that carries **both** the **SYN** and the **ACK** flags. To end this three-way handshake, the client sends a last TCP segment with the **ACK** flag, which acknowledges the receipt of the previous server response.
  - Data Transfer, where the data is encapsulated into TCP segments and is exchanged between the server and the client.
  - Connection Termination (Four-Way Handshake), where the client starts by sending a TCP segment with the **FIN** flag to the server, who responds with a segment carrying the **ACK** flag and another one that carries another **FIN** flag. Then, the client answers back with the last TCP segment, that carries an **ACK** flag.
4. How does the ARQ TCP mechanism work? What are the relevant TCP fields? What relevant information can be observed in the logs?  
Automatic Repeat Request (ARQ) is used to retransmit information on a congested network, which is a network where packets have been lost. In order to lose packets, many have to be sent at the same time (Additive Increase). We can use Slow Start when, instead of adding a one to CongestionWindow in each transmission, we double the CongestionWindow in each retransmission. Packet loss can either occur by timeout (there's a Multiplicative Decrease, making Congestion Window 1 and increasing it back until its half the value that's gotten via slow start. After that, Additive Increase starts being incremented by 1) or due to 3 consecutive ACKs (there's a Multiplicative Decrease, reducing Congestion Window to half and the Additive Increase starts being incremented by 1).
5. How does the TCP congestion control mechanism work? What are the relevant fields. How did the throughput of the data connection evolve along the time? Is it according to

the TCP congestion control mechanism?

Each sender determines the channel capacity in order to send one or less packets.

Because of that, there's another new parameter in the connection - CongestionWindow. If the network's congestion level increases, the CongestionWindow decreases and vice-versa.

6. Is the throughput of a TCP data connection disturbed by the appearance of a second TCP connection? How?

Yes - by creating more than once TCP connection, the bandwidth will be divided by all the connections, slowing each one of them down.

## Conclusions

By correctly implementing all the steps of the project, from the app to the experiences, we had the chance to better understand how the TCP protocol works, how the configuration of a complex network is done and how the many kinds of packets behave and what information they carry.

## Annexes

### 1. Code

- Available in folder app.

### 2. Wireshark Captures

- Logs and prints availables in resource folder.